



TECHNICAL UNIVERSITY OF DENMARK
DTU COMPUTE

CONCURRENT PROGRAMMING
FALL 2014

Car control

DUE NOVEMBER 19, 2014

Authors: Group 24

Michael BØNDERGAARD
(s113112)

Kristoffer BREITENSTEIN
(s113135)

Contents

1	Introduction	1
2	Step 1	2
3	Step 2	3
3.1	Analysis	3
3.2	Implementation	4
4	Step 3	4
4.1	Ideas & Design	4
4.2	Implementation	4
4.3	Extra B	5
5	Step 4	5
5.1	Discussion	5
5.2	Implementation	5
6	Step 5	5
6.1	Ideas & Design	5
6.2	Implementation	6
6.3	Extra G	6
7	Tests	6
8	Evaluation	6
9	Conclusion	6
10	Appendix	6
A	Collision avoidance	6
B	jSpin	7

1 Introduction

The course teaches how to handle software with additional tasks within one software program. Software will have threads handle different tasks and have shared variables, which the threads will have limited or full access to. Concurrent thread can be used for many different objects and goals depending on the kind of project the software is written for. Threads can often be used as optimization in software, where different threads will handle different part of an calculation or another task. Especially with the modern processors running multiple cores at once, then one core can handle a task while another handles a different task.

The project is about concurrent processes running individually and using specific data in cooperation with the other processes. The project has 9 cars driving around a parking lot with the cars having different routes. The cars will have to make sure they do not cause accidents or end up in a dead lock situation. The cars will have to pass through an alley with only room the cars driving in one direction. The cars are run by a thread each to make sure the cars are not getting into an accident.

In the project the alley is managed by either a semaphore or a monitor written in Java. These are two different approaches of how to handle atomic actions in software. In other words mean the variable or a critical section is only available to one thread at a time.

In the end the cars will be driving with the alley acting as a traffic light. The cars only being able to drive in one direction at a time, while the others will have to wait till the alley is available. The cars also has a barrier, where the cars will wait until a specific number of cars are waiting. The traffic will end up going around the circuit in a smooth flow.

2 Step 1

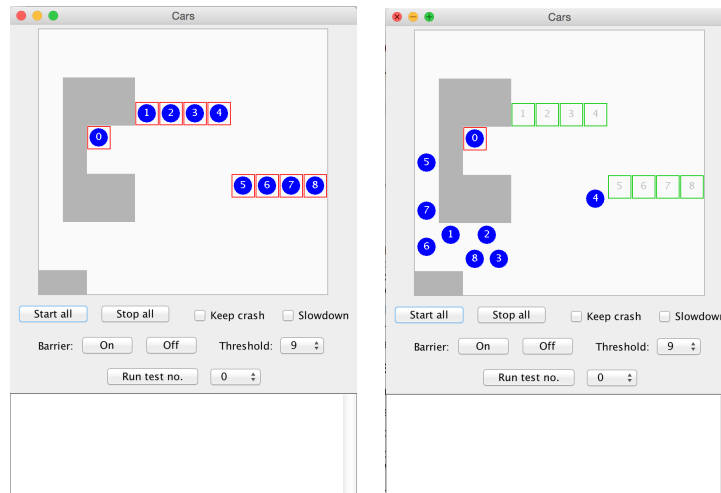
The system given has the cars driving around the circuit in their specific routes. The cars does not stop, when they bump into each other, instead a red square is shown for each collision.

Step one's requirements are to

- make sure the cars do not collide
- cause deadlocks in the traffic
- have cars drive through the alley

Overall the challenges is to stop the cars' collisions and have semaphores handle the alley. The deadlocks is a part in both the collision and the alley in a way.

Figure 1: The cars driving around the parking lot



The idea behind avoiding collision is that the cars can not move, if another car is either already in the way or both cars are trying to enter same field. The cars will in both cases stop. A car already being an obstacle, the waiting car will have to wait for the other car to move along. This situation will never cause a deadlock, but will have cars waiting for each other in a line.

The other case where two cars want to enter same field, can cause a deadlock with both cars waiting for the other car to move. A solution to the deadlock situation is to have a way to let one of the cars drive on some condition. The condition we have chosen is to let the cars with the highest number advance. The condition will have to be unique, because otherwise the cars will collide. The condition is unique, since no cars have the same number.

One problem with the solution is, when the cars are driving through the alley a direction will stop. The cars driving up the alley will stop for the cars with a higher number. In this case a solution is to let the cars in the alley being prioritized higher, than the cars waiting for the alley to be vacant.

3 Step 2

The software spin uses Promela models to verify a concurrent model. The model shows rather the processes uses variables and other shared features in a valid and intended way. The processes should not be allowed to enter more than 4 cars at a time and only in one direction.

3.1 Analysis

The Promela code is a little different to get to work than the java code. The semaphore described previously result in the code found in appendix B. The code display the enter part followed by the leave part of the cars alley class in java. The result of the analysis shows with 8 threads the program runs correctly.

```

pan: resizing hashtable to -w29.. pan: out of memory
hint: to reduce memory, recompile with
      -DCOLLAPSE # good, fast compression, or
      -DMA=84    # better/slower compression, or
      -DHC # hash-compaction, approximation
      -DBITSTATE # supertrace, approximation
(Spin Version 6.1.0 — 4 May 2011)
Warning: Search not completed
        + Partial Order Reduction
Full statespace search for:
      never claim           - (none specified)
      assertion violations  +
      cycle checks          - (disabled by -DSAFETY)
      invalid end states    +
State-vector 84 byte, depth reached 108, ... errors: 0 ...
  2.69e+08 states, stored
  6.5625499e+08 states, matched
  9.2525499e+08 transitions (= stored+matched)
  1.4695379e+08 atomic steps
hash conflicts: 4.0118364e+08 (resolved)
Stats on memory usage (in Megabytes):
28732.300      equivalent memory usage for
states (stored*(State-vector + overhead))
11478.569      actual memory usage for states
      (compression: 39.95%)
               state-vector as stored = 17 byte +
               28 byte overhead
  4096.000      memory used for hash table (-w29)
    0.107      memory used for DFS stack (-m2000)
    4.417      memory lost to fragmentation
15570.259      total actual memory usage
pan: elapsed time 1.49e+03 seconds
pan: rate      180388 states/second

```

The analysis does not show any errors, but a warning is given "Search not complete". Other analysis have been run, which does not show the warning,

but these are with less processes. These analysis can be found in the appendix B. The problem is that the pan runs out of memory on the thin clients, which result in a termination. The program does the analysis, but takes a long time and so far the biggest analysis run has 6 processes without error. The error can occur on 6 processes, but does not occur every time. For 5 processes the analysis does seem to work every time, but does only have 1 car in the one direction and 4 in the other.

The analysis shows no errors which means the checks (found at the bottom of the code) do hold. The different semaphores are kept within the limits of the variables.

3.2 Implementation

The code in appendix B is a close implementation of the code in java, which can be found in appendix ??

Some of the challenges has been to get the jSpin code to do the exact same as the java code, but we have received some help from the TA to get it right. The challenges was to make the wait function along with the semaphore code to behave as intended.

4 Step 3

This step concerns the barrier. The barrier has to stop the cars until all cars are waiting together. Then they should all be released for a round, and wait when they arrive at the barrier next time.

4.1 Ideas & Design

Before and during the implementation, some ideas were discussed. The cars had to stop at the barrier. Then they had to be released, but only until they arrived at the barrier again. Also the cars should stop waiting, when the barrier is turned off.

The idea was to make the on and off function quite similar to the on and off function from the Gate-class. The only difference it that the off function had to release the waiting cars.

The hard problem was, how to do the sync function right. A counter is used to count waiting cars, and a single semaphore is used to make the cars wait.

The sync function should keep track of number of cars, and whether the barrier is on or off.

4.2 Implementation

The on and off functions were implemented as the on and off functions from the Gate-class almost. The difference is that the off function releases the cars waiting at the barrier. The on function takes the only coconut in the semaphore b.

As seen, the sync function increases the cars counter, checks on the barrierOn boolean. Nothing happens in the case that the barrier is not on, meaning off.

The cars counter will just be decreased again and the car will go on. In the case that the barrier is on, the car will wait for a coconut in the semaphore, b. Until the cars counter reaches 9, the cars will wait. The semaphore b had only a single coconut, which was taken by the on function. When the 9th car comes along, the cars counter will not be less than 9, and it will release a number of coconuts corresponding to the cars counter. This means, that all cars will be released, and the

4.3 Extra B

5 Step 4

The task of step 4 is to implement the alley and the barrier by monitors instead of semaphores.

When using monitors the semaphore class is not needed. The functions are extended to be synchronized functions, and the thread calls "notifyAll" and "wait" is used.

5.1 Discussion

It is more convenient to use monitors for the barrier and the alley. Since the alley kind of uses groups and the barrier too. The monitor is more convenient having group whereas the semaphore is more oriented against only letting one thread pass at the time.

By using monitors the code is also much less complicated which will be shown during this section.

5.2 Implementation

6 Step 5

This step is about how to remove a car and restore it again, not necessarily immediately after. Subproblems of this task is, how to stop the car, how to remove the car, and how to restore the car. A even harder question is, what to do when removing a car inside the alley or waiting at the barrier. These questions are considered during this task.

6.1 Ideas & Design

- How to stop the car?
- How to remove the car?
- How to deal with the remove of the car for the alley and barrier?

6.2 Implementation

6.3 Extra G

7 Tests

The testing of the program has been made during the programming phase. But also some tests have been made in the CarTest-class to show examples of the program running.

8 Evaluation

9 Conclusion

10 Appendix

A Collision avoidance

The code for the cars to avoid collision can be seen below.

```
for (int i = 0; i < 9; i++) {
    if (i != no) {
        //Checking the position of the other cars
        //with the new position
        if (position[i].equals(newpos)) {
            free = false;
        }
        //checking the next position of the other
        //cars with the cars next position
        if (newpos.equals(cd.nextPos(i, position[i]))) {
            if(inAlley){
                if(no<5){
                    //Make sure the cars in
                    //the alley goes first
                    if (no > i) {
                        free = false;
                    }
                }else{
                    if (no < i) {
                        free = false;
                    }
                }
            }
        }else{
            //If the cars are outside the alley the car
            // with highest number goes first
            if (no < i) {
                free = false;
            }
        }
    }
}
```



```

    }
}

```

B jSpin

The code from jSpin to run the analysis:

```

define N                8          /* no. of processes */

short c = 0;
short u = 4;
short d = 1;
short b = 1;
bool trafficUp = false;

/* Declare and instantiate N Counter processes */

inline v(s){
    s++;
}

inline p(s){
    atomic{
        s > 0 -> s--;
    }
}

active [N] proctype Alley ()
{
    c=(c+1)%N;
entry:
    if :: c<4 ->
        if :: trafficUp == true ->
            if :: u == 4 -> p(d);
            trafficUp = true;
        :: else -> skip;
        fi;
    p(u);
    :: trafficUp == false -> p(b);
        if :: trafficUp == true ->
            v(b);
        :: trafficUp == false ->
            p(d);
            v(b);
        fi;
    p(u);
    trafficUp=true;
    :: else -> skip;
    fi;
}

```

```

:: c>=4 ->
    if :: trafficUp == true ->
        p(d);
        p(u);
        trafficUp = false;
    :: trafficUp == false; ->
        if :: u == 4 ->
            p(d);
            trafficUp=false;
        :: else -> skip;
        fi;
        p(u);
    fi;
:: else -> skip;
fi;

leave:
    v(u);
    atomic{
        if :: u==4 ->
            if :: d== 0 -> v(d);
            :: else -> skip;
            fi;
        :: else -> skip;
        fi;
    }
}

/* Invariant check */
active proctype Check ()
{
    (0 <= c && c <= 7) -> assert(true);
    (0 <= u && u <= 4) -> assert(true);
    (0 <= d && d <= 1) -> assert(true);
    (0 <= b && b <= 1) -> assert(true);
}

```

The analysis run with 5 processes:

(Spin Version 6.1.0 — 4 May 2011)
+ Partial Order Reduction

Full statespace search for:

never claim	- (none specified)
assertion violations	+
cycle checks	- (disabled by -DSAFETY)
invalid end states	+

State-vector 60 byte, depth reached 72, ... errors: 0 ...

36981340 states, stored

57403474 states, matched

94384814 transitions (= stored+matched)

10018950 atomic steps

```

hash conflicts: 35003377 (resolved)
Stats on memory usage (in Megabytes):
 3103.598      equivalent memory usage for states
                (stored*(State-vector + overhead))
 2015.764      actual memory usage for states (compression: 64.95%)
                state-vector as stored = 29 byte + 28 byte overhead
    256.000    memory used for hash table (-w25)
    0.107      memory used for DFS stack (-m2000)
 2271.064      total actual memory usage
unreached in proctype Alley
                (0 of 91 states)
unreached in proctype Check
                (0 of 9 states)
pan: elapsed time 145 seconds
pan: rate 255784.62 states/second

The analysis run with 6 process:

pan: resizing hashtable to -w29.. pan: out of memory
hint: to reduce memory, recompile with
      -DCOLLAPSE # good, fast compression, or
      -DMA=68   # better/slower compression, or
      -DHC # hash-compaction, approximation
      -DBITSTATE # supertrace, approximation
(Spin Version 6.1.0 — 4 May 2011)
Warning: Search not completed
        + Partial Order Reduction
Full statespace search for:
      never claim                - (none specified)
      assertion violations      +
      cycle checks              - (disabled by -DSAFETY)
      invalid end states        +
State-vector 68 byte, depth reached 84, ... errors: 0 ...
 2.69e+08 states, stored
 4.5920728e+08 states, matched
 7.2820728e+08 transitions (= stored+matched)
 79882859 atomic steps
hash conflicts: 2.7405665e+08 (resolved)
Stats on memory usage (in Megabytes):
24627.686      equivalent memory usage for
states (stored*(State-vector + overhead))
11490.506      actual memory usage for states (compression: 46.66%)
                state-vector as stored = 17 byte + 28 byte overhead
    4096.000    memory used for hash table (-w29)
    0.107      memory used for DFS stack (-m2000)
    4.537      memory lost to fragmentation
15582.076      total actual memory usage
pan: elapsed time 1.16e+03 seconds
pan: rate 231499.41 states/second

```